

# Module 1: Getting Familiar with RC




1


**Be Boulder.**

# Learning Goals

1. Understand CURC Resources & the Alpine cluster
2. Getting an account & logging in
3. Navigate the RC system




**Ask Questions**



**Discuss Ideas**

---



Research Computing  
UNIVERSITY OF COLORADO BOULDER

3

**Be Boulder.**

This workshop is meant to be empowering and information-rich – to help you get the most out of your HPC workflow.

Feel free to ask questions during the presentations and to discuss ideas/thoughts during breaks and the hands-on portion of the workshop.

# RC Resources

## Tech Support

- High Performance Computing
- Data Management
- Cloud Computing
- Secure Research

## Human Support

- Training Materials & Workshops
- Consultations & Office Hours
- Help Tickets




4

**Be Boulder.**

The Secure Research manages the CMMC environment (Preserve), but we can help users connect with the SR team if they have confidential data (e.g. HIPA, Financial, FERPA, Government data, etc.)

# High Performance Computing (HPC)



# Traditional vs HPC



**Fewer Computing Resources,  
Fast Sprinter\***



**More Computing Resources,  
Steady Marathon Runner\***


## Traditional PC:

- 1 CPU with a few cores and maybe a GPU (not always)
- RAM & Storage measured in GB's
- Less Hardware, but potentially much faster hardware (clockspeed)


## HPC

- Many CPU's and GPUs – Thousands of cores
- RAM & Storage measured in Terabytes and Petabytes
- Hardware built for stability and high up-time, best for distributed jobs

# What can I use HPC for?



**Parallel Jobs**



**Serial Jobs**


Parallel Jobs - Divide and Conquer

- A job that can be divided into smaller pieces, which can be independently run.


Serial Jobs

- A collection of small jobs which can be done simultaneously

## What can I use HPC for?



**Big Data**



**Viz/Rendering**



**Shared Work  
Environment**

### Big Data:

- Require lots of RAM (memory) or data storage
- Need more CPU's, distributed system

### Visualization Rendering

### Unique Hardware

- GPU's, Linux, Shared team environment



# HPC Cluster: Alpine



- 3rd-generation HPC cluster
  - Janus (2012)
  - RMACC Summit (2017)
- Heterogeneous cluster with hardware currently provided by CU Boulder, CSU, and Anschutz
- Access available to CU Boulder, CSU, AMC and RMACC users

Update Icon

# Alpine Partitions

Partition	Description	# of nodes	RAM/core (GB)	cores/node	GPUs/node
amilan	General Compute Node: AMD Milan	<u>347</u>	3.74	64	0
ami100	GPU Node: 3x AMD MI100	8	3.74	64	<u>3</u>
aa100	GPU Node: 3x Nvidia A100	12	3.74	64	<u>3</u>
amem	High-memory node	22	<u>21.5</u>	64 (10), 48 (12)	0

# HPC Cluster: Alpine



- Interconnect
  - **CPU nodes:** HDR-100 InfiniBand\* (200Gb inter-node fabric)
  - **GPU nodes:** 2x25 Gb Ethernet +RoCE
  - **Scratch Storage:** 25Gb Ethernet +RoCE
- Operating System
  - RHEL 8 - RedHat Enterprise Linux Version 8




\* Not all CPU Nodes have Infiniband


## Node Types



**LOGIN**



**COMPILE**



**COMPUTE**

Add diagram showing login node as entry point.

Login -> Testing, Compiling, Compute/GPU

## Node Types



### LOGIN

- Entry to system
- View or edit files
- Submit jobs



### COMPILE



### COMPUTE


Add diagram showing login node as entry point.

Login -> Testing, Compiling, Compute/GPU

## Node Types




LOGIN



COMPILE

- View or edit files
- Submit jobs
- **Compile code**
- **Install software**



COMPUTE

Add diagram showing login node as entry point.

Login -> Testing, Compiling, Compute/GPU


## Node Types



LOGIN



COMPILE




COMPUTE

- Run scheduled jobs
- Handle calculations

Add diagram showing login node as entry point.

Login -> Testing, Compiling, Compute/GPU

## Where do I run jobs?



Add diagram showing login node as entry point.

Login -> Testing, Compiling, Compute/GPU



# Node Types



## DATA (DTN)

- Support data transfers
- Can be selected when using scp, sftp, or ssh transfers

# Storage at CURC



## Core

Included with RC account

- /home
- /projects
- /scratch



## PetaLibrary

Paid Service for large data

- Storage
- Archival Backup
- Sharing

# Core Filesystem Structure

**/home (2GB)**

**/projects (250GB)**

**/scratch/alpine (10TB)**

- Scripts and Code
- Important files (small)

**Not for:**

- File Sharing
- Large Files

Configuration Files,  
Notes



# Core Filesystem Structure

**/home (2GB)**

- Scripts and Code
- Important files (small)

**/projects (250GB)**

- Code/files/libraries
- Software you are installing
- Sharing files

**/scratch/alpine (10TB)**



**Not for:**

- File Sharing
- Large Files

**Not for**

- Job Output
- Temporary Files

Configuration Files,  
Notes

Job Scripts,  
Shareable Files

# Core Filesystem Structure

/home (2GB)	/projects (250GB)	/scratch/alpine (10TB)
<ul style="list-style-type: none"> <li>• Scripts and Code</li> <li>• Important files (small)</li> </ul>	<ul style="list-style-type: none"> <li>• Code/files/libraries</li> <li>• Software you are installing</li> <li>• Sharing files</li> </ul>	<ul style="list-style-type: none"> <li>• Output from running jobs</li> <li>• Large files/datasets</li> <li>• Sharing files</li> <li>• Cluster specific</li> </ul>
<b>Not for:</b> <ul style="list-style-type: none"> <li>• File Sharing</li> <li>• Large Files</li> </ul>	<b>Not for</b> <ul style="list-style-type: none"> <li>• Job Output</li> <li>• Temporary Files</li> </ul>	<b>Not for</b> <ul style="list-style-type: none"> <li>• Long Term Storage</li> </ul>
Configuration Files, Notes	Job Scripts, Shareable Files	Job Output Files, Data Files



# PetaLibrary Tiers

Active	Archive	Active + Archive	Archive + DR
<ul style="list-style-type: none"> <li>• Performance Tier<sup>1</sup></li> <li>• Accessible by All Nodes</li> <li>• No File Limit</li> <li>• Double Parity</li> <li>• lz4 compression</li> </ul>			

<sup>1</sup> Slower than /scratch, <sup>2</sup> Per Terabyte

# PetaLibrary Tiers

Active	Archive	Active + Archive	Archive + DR
<ul style="list-style-type: none"> <li>• Performance Tier<sup>1</sup></li> <li>• Accessible by All Nodes</li> <li>• No File Limit</li> <li>• Double Parity</li> <li>• lz4 compression</li> </ul>	<ul style="list-style-type: none"> <li>• Integrity Tier</li> <li>• Accessible by Login Nodes Only</li> <li>• 10,000 File Limit<sup>2</sup></li> <li>• Triple Parity</li> <li>• zstd compression</li> </ul>		

<sup>1</sup> Slower than /scratch, <sup>2</sup> Per Terabyte

# PetaLibrary Tiers

Active	Archive	Active + Archive	Archive + DR
<ul style="list-style-type: none"> <li>• Performance Tier<sup>1</sup></li> <li>• Accessible by All Nodes</li> <li>• No File Limit</li> <li>• Double Parity</li> <li>• lz4 compression</li> </ul>	<ul style="list-style-type: none"> <li>• Integrity Tier</li> <li>• Accessible by Login Nodes Only</li> <li>• 10,000 File Limit<sup>2</sup></li> <li>• Triple Parity</li> <li>• zstd compression</li> </ul>	<ul style="list-style-type: none"> <li>• Replicate Active into Archive</li> <li>• Updated every 15 minutes</li> <li>• No File Limit</li> </ul>	

<sup>1</sup> Slower than /scratch, <sup>2</sup> Per Terabyte




# PetaLibrary Tiers

Active	Archive	Active + Archive	Archive + DR
<ul style="list-style-type: none"> <li>• Performance Tier<sup>1</sup></li> <li>• Accessible by All Nodes</li> <li>• No File Limit</li> <li>• Double Parity</li> <li>• lz4 compression</li> </ul>	<ul style="list-style-type: none"> <li>• Integrity Tier</li> <li>• Accessible by Login Nodes Only</li> <li>• 10,000 File Limit<sup>2</sup></li> <li>• Triple Parity</li> <li>• zstd compression</li> </ul>	<ul style="list-style-type: none"> <li>• Replicate Active into Archive</li> <li>• Updated every 15 minutes</li> <li>• No File Limit</li> </ul>	<ul style="list-style-type: none"> <li>• Disaster Recovery</li> <li>• Monthly backup to offsite location</li> </ul>

<sup>1</sup> Slower than /scratch, <sup>2</sup> Per Terabyte

## Data Sharing: Within RC




\*Users manage access to their **/project** and **/scratch** directories.

RC Staff manage access to the PetaLibrary Space

# Data Sharing: Outside RC

## Large Data Transfers:

- Globus (Recommended)
- Data Transfer Nodes (DTN)
- Terminal/Command Line:
  - `rsync`
  - `rclone`
  - `sftp`
  - `scp`



DTN require internal CU Network access (VPN)

## Acceptable data storage and use

CURC systems and services **should not be used to store** any data that is US government Classified, nor any Controlled Unclassified Information.

For users requiring storage for sensitive data types, please see the secure research computing resources:

<https://www.colorado.edu/rc/secure-research-computing-resources>

# Cloud Computing

- CURC supports AWS, Azure, and GCP
- Alternative to HPC
- Enhance HPC



Cloudbursting

# How to Access RC Resources?

1. Get an RC account
2. Set up two-factor authentication with Duo
3. Log in
4. Create greatness! (responsibly)

# Getting an Account

- CU Boulder, CSU users and affiliates:
  - Request an account through the RC Account request portal:  
<https://rcamp.rc.colorado.edu/accounts/account-request/create/organization>
- AMC, RMACC users and affiliates:
  - Request an account through the ACCESS-CI User Registration Portal:  
<https://identity.access-ci.org/new-user.html>

# Your RC Account

## Access to:

1. Alpine Cluster
2. Core Storage
3. PetaLibrary Storage\*
4. Open OnDemand
5. Approximately 2,000 Service Units (SUs) per month

\*If purchased



32


**Be Boulder.**

<https://ondemand.rc.colorado.edu/>  
<https://ondemand-rmacc.rc.colorado.edu/>



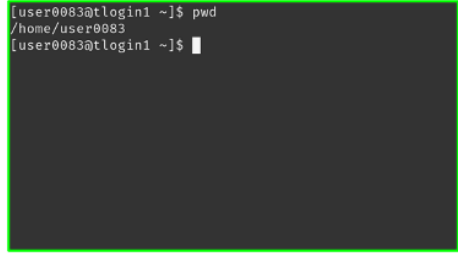
# Duo Authentication

1. Duo smartphone app (recommended)
2. Phone Call/Text is an alternatives



# Terminal Access

- Mac or Linux
  - Terminal application
- Windows
  - PuTTY
  - Powershell
- Open OnDemand (*alternative for CU affiliates*)
  - For those less familiar with Linux ([ondemand.rc.colorado.edu/](https://ondemand.rc.colorado.edu/))

A screenshot of a terminal window with a dark background and green text. The prompt is [user0083@tlogin1 ~]\$. The command 'pwd' has been entered and executed, returning the path /home/user0083. The prompt is now [user0083@tlogin1 ~]\$ followed by a cursor.

```
[user0083@tlogin1 ~]$ pwd
/home/user0083
[user0083@tlogin1 ~]$
```

# Demo: Logging in via Terminal

To login to an RC login node:

```
$ ssh <username>@login.rc.colorado.edu
```

Supply your IdentiKey\* password and your Duo app will alert you to confirm the login

\* Exclusive to CU and CSU accounts

# Demo: logging in with OnDemand

CURC Open OnDemand is a browser based, integrated, single access point for all of your HPC resources at CU Research Computing.

- CU Boulder: Visit <https://ondemand.rc.colorado.edu>.
- Other RMACC Institutions: Visit <https://ondemand-rmacc.rc.colorado.edu/>